LOFAR User Guide:
The Framework Pulsar Pipeline

Pulp, Version 1.0

Document version 1.0

SVN Repository .Revision: 6993

K. R. Anderson

SVN Date: 2011-02-05

Contents

6 Appendix A 21

Change record

Issue	Date	Sections	Description of changes
1.0	2011-01-31	all	Initial release

1 Introduction

`>>> import` This document is version 1.0 of the LOFAR User Guide for the LOFAR framework pulsar pipeline, known colloquially as 'pulp'. This User Guide provides general instruction to help users establish a viable user environment as required by the LOFAR pipeline framework in order to run the framework pulsar pipeline in the LOFAR cluster compute domain.

This document will also present details on running the framework pulsar pipeline, known as Pulp (executable: `pulp.py`). Users are encouraged to seek detailed information about the actual processing engaged by the "known pulsar pipeline" from the LOFAR Transient Key Project (TKP) science team. Details regarding LOFAR beam-formed data processing actual are beyond the scope of this document.

This document assumes some familiarity with terms surrounding the pipeline framework, such as "recipe", and "head node", as well as terms associated with "known pulsar pipeline" processing. Such terms will be used throughout this document. [1]

1.1 Applicable Documents

Documentation on the Pulp package is provided with the Pulp distribution, and is available with the Pulp package download in `${LOFARSOFT}/src/Pulsar/pipeline/documentation`[2]. Further documentation on the LOFAR pipeline framework, development under the LOFAR pipeline framework, as well as schematic documentation on the Pulp API, are available within the repository itself. Pulp API documents are delivered with the Pulp package as part of any download of the `$LOFARSOFT` repository.

The Pulp current release under the LOFAR USG repository is located in `$LOFARSOFT/src/Pulsar/pipeline`:

```
drwxr-xr-x   9 <user>    staff      306 Feb  2 18:07 .svn
-rw-r--r--@  1 <user>    staff       67 Jan 31 01:29 __init__.py
drwxr-xr-x   8 <user>    staff      272 Feb  2 18:05 documentation/
-rw-r--r--   1 <user>    staff     3614 Jan 28 20:33 dynspec.py
-rw-r--r--   1 <user>    staff     1708 Jan 21 17:16 pipeline.cfg
-rwxr-xr-x   1 <user>    staff     4545 Jan 31 01:48 pulp.py
rw-r--r--    1 <user>    staff       40 Jan 28 17:28 pulpVersion.py
drwxr-xr-x   6 <user>    staff      204 Jan 31 01:21 recipes/
drwxr-xr-x  18 <user>    staff      612 Feb  1 16:16 support/
-rw-r--r--   1 <user>    staff      789 Jan 28 22:07 tasks.cfg
```

In `documentation/`, users will find a directory structure that somewhat mimics the pipeline's recipe-support directory structure, with recipe API documents found under `./pipeline/documentation/PulpAPI/` in 'master','nodes', and 'support', which contain, respectively, head node recipes, compute node recipes, and support modules providing adminstrative and executive function.

`$LOFARSOFT/src/Pulsar/pipeline/documentation`:

```
  drwxr-xr-x   9 <user>     staff  306 Jan 30 08:08 .svn
  drwxr-xr-x   6 <user>     staff  204 Feb  3 15:09 PresentationsAndPapers/
  drwxr-xr-x   6 <user>     staff  204 Feb  3 15:09 PulpAPI/
  drwxr-xr-x  27 <user>     staff  918 Feb  3 00:00 UserGuide/
```

where a user will find API documents for the pipeline recipes, and all support modules.

```
./PulpAPI/master:
-------
-rw-r--r--@   0  __init__.py
 rw-r--r--  ...  bf2presto.html
```

---

[1] Readers of this document unfamiliar with the terms, 'known pulsar pipeline,' or 'Jason Hessels', should probably be reading something else.

[2] This document can be found op "UserGuide/pulpUserGuide.pdf'," under "documentation/"

```
-rw-r--r--  ...   buildPulsArch.html
-rw-r--r--  ...   buildRSPAll.html
-rw-r--r--  ...   bundleFiles.html
-rw-r--r--  ...   prepareInf.html
-rw-r--r--  ...   prepfold.html
-rw-r--r--  ...   rfiplot.html

./PulpAPI/nodes:
------
-rw-r--r--@   0  __init__.py
-rw-r--r--  ...   bf2presto.html
-rw-r--r--  ...   prepfold.html
-rw-r--r--  ...   rfiplot.html

./PulpAPI/support:
--------
-rw-r--r--@   0  __init__.py
 rw-r--r--  ...   RSPlist.html
-rw-r--r--  ...   bf2Pars.html
-rw-r--r--  ...   buildRSPS.html
-rw-r--r--  ...   bundlePlots.html
-rw-r--r--  ...   foldingData.html
-rw-r--r--  ...   fullRSP.html
-rw-r--r--  ...   pardata.html
-rw-r--r--  ...   prepInfFiles.html
-rw-r--r--  ...   pulpEnv.html
-rw-r--r--  ...   rfiDirectories.html
```

These are accessible and viewable through any browser[3]

1.2 Reference Documents

For documentation on the framework itself, as well as guidance on recipe writing, the user is invited to examine the framework documenation available in the repository under

```
${LOFARSOFT}/src/pipeline/docs .
```

A subset of the framework's API is documented in the Pulp API documentation described above.

1.3 Package Overview

The Pulp package is a framework implementation of the "basic mode" of the LOFAR TKP Pulsar Group's pipeline shell script[4]. Pulp is built on the LOFAR pipeline framework and, as with all such pipelines, comprises a set of 'recipes'[5]that are executed in successive order. In a cluster compute environment, these recipes can be paired across push-pull, head-compute node connections in order to facilitate parallel, multi-compute-node processing on a cluster and/or subsections of a cluster. Depending upon the required action, recipes may or may not be paired across head node/compute node connections.[6] When a "head node" recipe is not paired with a matching "compute node" recipe (name matching of head/compute node recipes is required by the framework), this necessarily implies that a single process will be activated on a compute

---

[3]Firefox (v3.0+) has a nice indexing feature and is recommended for viewing these documents.

[4]make_subs_SAS_Ncore_Mmodes.sh, src: ${LOFARSOFT}/lofarsoft/src/Pulsar/scripts

[5]The term "recipe" is apprently being used in other developing and developed pipeline frameworks; noted at ADASS XX, Boston (2010).

[6]This is explicitly true for the Pulp package, as it employs the framework's IPython facility. This statement is likely not correct for the "new style" imaging pipeline, which employs an ssh protocol.

node in order to perform a task that does not require parallelization. For example, building the storage node directory structure for an observation.

In order to illustrate this meaning, Table 1 presents the overall organisation of the Pulp pulsar pipeline package. The table demonstrates those recipes that will perform the tasks requiring multiple parallel jobs. Those recipes (bolded) show matching named "node recipe" modules in the package. It is these head node/compute node recipe pairs that multiply execute all requested processing jobs. Those jobs, and the resulting job queues, are arranged and built by the respective head node recipe, and then farmed to the cluster based upon a user's "clusterdesc" file (see §2).

MASTER RECIPE	NODE RECIPE	SUPPORT MODULES
buildPulsArch	—	buildRSPS,RSPlist
bf2presto	**bf2presto**	bf2Pars, pulpEnv
buildRSPALL	—	fullRSP, pulpEnv
prepareInf	—	prepareInfFiles, pulpEnv
prepfold	**prepfold**	pulpEnv
rfiplot	**rfiplot**	rfiDirectories, pulpEnv
bundleFiles	—	bundlePlots, pulpEnv

Table 1: Pulp distribution and recipe relations

Note: `support/` modules may support other `support/` modules. Recipes do not support other recipes. As indicated earlier (§1), the described Pulp package is delivered via checkout of the LOFAR USG code repository.

1.3.1 Nominal Pulp Processing

The nominal processing performed by the Pulp package involves (terms used here are explained in detail in §2, and §3),

- The number of subbands in a observation is entirely selected by the observer, however, nominally 248 subbands of 'incoherent' beam-formed data are usually delivered and processed. Currently, Pulp can only process so-called 'incoherentstokes' beam-formed data.[7][8]

- The pipeline is capable of handling any user selected "splitting" factor, i.e. the 'filefactor' parameter within Pulp. This factor is known as "ncores" in the pulsar shell script[9]. Pulp places no restriction on what 'filefactor' can be, though nominal processing will usually dictate the default `filefactor = 8`. Because of this, users are encouraged to use judgment in filefactor specification, when the default is not desired.

- Pulp accomodates any number of observation subbands and any splitting factor, though the number of subbands appears unlikely to exceed 248. Nominal beam-formed data delivery is expected to range from a few files[10], to the usual 248. However, Pulp places no constraint on the number of subbands/files an observation may comprise[11].

- In the special case that a user select `filefactor == 1`, No "all" processing will be done. That is, no `RSPA` is made, and no `RSPA` processing, *per se*, occurs. Output will appear in the expected 'RSP0' directory. RSP0 is equivalent to RSPA, when `filefactor=1`.

---

[7]Development on 'coherentstokes' data processing tabled until beam-formed data type is written in conformance with the LOFAR-USG-ICD-003 specification.

[8]see LOFAR document, *LOFAR-USG-ICD-003, Beam-formed Data*, Alexov et al, ASTRON, 2010.

[9]Pulp is node generic at the user interface and the API.

[10]esp. "coherentstokes" beam-formed data.

[11]Pulp seemlessly handles "uneven" splitting of subbands across RSP groups. In cases where evaluation of the quantity `nsubbands mod filefactor` is uneven divison, the modulo number of subbands are added to the "last" RSP group, i.e. RSP[filefactor - 1].

- The "rfi report" is produced automatically and is not currently switchable[12]. Pulp executes the relevent steps automatically and the products, a dynamic spectrum of the data and a so-called ".rffireport" file, are bundled as a standard pipeline data product.

Details on the Pulp command line interface are provided in §3, "Interfaces".

1.4   Glossary

- **API** – Application programming interface (!Anton Pannekoek Instituut).

- **Compute node** – One of an assigned number of cluster nodes configured for computing. For LOFAR, these compute nodes have NFS connections to selected storage nodes, which is where both input and output data will be read from and written to[13]. Within the context of the LOFAR pipeline framework, compute nodes are often referred to simply as "nodes".

- **Cluster** – A multi-node computing cluster.

- **Head node** – The head (controlling) node of a cluster environment. In the LOFAR computing environment, the head node is lfe001[14].

- **Pulp** – the "known pulsar pipeline" implemented under the LOFAR pipeline framework.

- **TKP**– The Transient Key Project (LOFAR).

2   User Environment

In computing environments, pipelines function as automated, simulated users at the command line. More aptly, a pipeline functions as a proxy user, or user agent. Therefore, and in order that the Pulp package execute properly, Pulp users will need to take a number of steps in adjusting their compute environments. Framework pipeline operations require that a number of locally defined configuration files be accessible to the framework. These files are (in no particular order):

1. pipeline.cfg – configuration for framework operations

2. tasks.cfg – define things to do

3. sub[n].clusterdesc – define target compute nodes

4. task.furl – job control

5. multiengine.furl – ip engine control

The two configuration files, "pipeline.cfg" and "tasks.cfg," listed are provided with the Pulp download. Described below, users will want to edit the pipeline configuration files to use their own builds of the `$LOFARSOFT` code repository. Users may edit the `tasks.cfg` file, which is where a user may easily change default arguments of any or all recipes. In all likelihood, most users will not want or need to do this. These configuration files, and others, are discussed in detail in §2.1.2, "Framework configuration files".

---

[12]Switching "rfireport" production (or any other part of the pipeline) is as easy to do as providing a command line switch, as the recipe(s) can simply be excluded or included from a user's pipeline definition with a single comment key. And the user will have developed their own customized user interface to Pulp.

[13]IO bottleneck

[14]Rumours of a second "head node" on the LOFAR cluster, "lfe002", have been reported, though remain unconfirmed by this user.

2.1 Configuration

2.1.1 Environment variables, paths

The following steps should be taken by users to establish the environment variables and resource `PATH`s required by the pipeline framework.

- – Define `LOFARSOFT`, `LOFARROOT` (`=/opt/LofIm/daily/lofar`)

- – source `${LOFARSOFT}/devel_common/scripts/init.sh`

- – Use `LofIm`

- – Adjust `$PYTHONPATH` to include framework python libraries.

Invocation of '`Use LofIm`' will define the environment variables, `$TEMPO`, and `$PRESTO`, needed for embedded PRESTO[15] operations.

Once the paths are defined (usually within a user's shell resource file), a user's `$PYTHONPATH` should further include the following paths described below. This should become obsolete in the near future (v1.1), with refactoring to use package namespace explicivity, rather than some still "blind" internal imports[16]

```
PYTHONPATH=/opt/pipeline/dependencies/lib/python2.5/site-packages:\
/opt/pipeline/framework/lib/python2.5/site-packages:               \
/opt/LofIm/daily/pyrap/lib:                                        \
${LOFARROOT}/lib/python2.5/site-packages:                         \
/opt/pythonlibs/lib/python/site-packages:                         \
${LOFARSOFT}/src/Pulsar/pipeline/recipes/master:                  \
${LOFARSOFT}/src/Pulsar/pipeline/recipes/nodes:                   \
${LOFARSOFT}/src/Pulsar/pipeline/support
```

A user's `$PATH` environment variable should be adjusted to include the following paths.
`$PATH` include:

```
${LOFARSOFT}/release/bin:\
${LOFARSOFT}/release/share/pulsar/bin:\
${LOFARSOFT}/src/Pulsar/pipeline:\
/opt/LofIm/daily/casarest/bin:\
/opt/pipeline/dependencies/bin:\
/opt/LofIm/daily/askapsoft/bin:\
/opt/scripts:${PATH}
```

N.B. These paths are subject to future alteration of the LOFAR system.[17]

2.1.2 Framework configuration files, directories

The pipeline framework functions through the use and availability of a limited set of configuration files, i.e., `.cfg` files. These configuration files are listed and briefly described. For more information, users are encourage to consult the documentation on the pipeline framework (Swinbank, 2011).

```
pipeline.cfg
tasks.cfg
```

---

[15]http://www.cv.nrao.edu/~sransom/presto
[16]on `support/` modules.
[17]Scrubbing of unused imaging pipeline imports has been performed to a large extent. However, clutter may still appear, eg., the `casarest` and `pyrap` paths from this user's own path sample, which are "this" environment necessary.

Users will notice that these two configuration files are delivered by the svn repository. These files are required by the framework. They are *not* part of the Pulp package. Though they must be tailored by the user to their particular environment and file system, the files are provided by the repository as both convenience and necessity; the `pipeline.cfg` the convenience, the `tasks.cfg` the necessity.

Users will need to make a directory the framework expects to use for pipeline activities, such as writing various log files. Within the above mentioned configuration file, '`pipeline.cfg`', detail below, users will define the configuration parameter, '`runtime_directory`'.

i.e.

```
runtime_directory = /path/to/your/runtime_directory
```

As illustrated, this directory is arbitrary in both name and location, though in nominal practice, it is usually located under a user's home directory, while "`runtime_directory`" is also somewhat conventionalized.[18] The following steps should be taken by users to establish a configuration required and recognized by the pipeline framework.

 – Make a "`pipeline_runtime/`" directory (arbitrary location)

 – Edit user's `pipeline.cfg` file to reflect this path

The `pipeline.cfg` file is a "pipeline generic" configuration file applicable to all framework pipelines. The Pulp `tasks.cfg` file is not. It is advised that this file be altered only with perspicacious intent.

Here is the top level of a delivered `$LOFARSOFT` Pulp directory:

```
[...]/lofarsoft/src/Pulsar/pipeline:
  drwxr-xr-x      .
  drwxr-xr-x      ..
  drwxr-xr-x      .svn
  -rw-r--r--@ ... __init__.py
  drwxr-xr-x  ... documentation/
  -rw-r--r--  ... dynspec.py
  -rw-r--r--  ... pipeline.cfg  ---> user paths and locations
  -rwxr-xr-x  ... pulp.py
  -rw-r--r--  ... pulpVersion.py
  drwxr-xr-x  ... recipes/
  drwxr-xr-x  ... support/
  -rw-r--r--  ... tasks.cfg     ---> task/recipe definitions
```

The reasons these particular files are delivered as part of the Pulp package are two fold: provide user's with a configuration basis, and to specify Pulp functionality, delivered by the `tasks.cfg` file. This confiuration file is critical to proper pipeline operations and should be altered only with user percipience of pipeline and framework operations.

Here is a quasi-generic pipeliine configuration file[19]:

```
[DEFAULT]
runtime_directory  = /some/path/pipeline_runtime
recipe_directories = [<LOFARSOFTpath>/src/Pulsar/pipeline/recipes]
lofarroot          = /opt/LofIm/daily/lofar
default_working_directory = /data/scratch/<user>
task_files         = [<LOFARSOFTpath>/src/Pulsar/pipeline/tasks.cfg]


[layout]
job_directory      = %(runtime_directory)s/jobs/%(job_name)s
```

[18]Explicit paths must appear in this configuration file.

[19]For details on the various sections of the `tasks.cfg` file, users should consult the documentation on the frameword pipeline (Swinbank, 2011), indicated §1.2

```
log_directory     = %(job_directory)s/logs/
vds_directory     = %(job_directory)s/vds
parset_directory  = %(job_directory)s/parsets
results_directory = %(job_directory)s/results/%(start_time)s


[cluster]
clustername      = pulsar
#clusterdesc     = %(runtime_directory)s/sub5.clusterdesc
clusterdesc      = %(runtime_directorypipe)s/line_runtime/sub6.clusterdesc
task_furl        = %(runtime_directory)s/task.furl
multiengine_furl = %(runtime_directory)s/multiengine.furl


[deploy]
script_path      = /opt/pipeline/framework/bin
controller_ppath = <LOFARSOFTpath>/src/Pulsar/pipeline/support:\
/opt/pipeline/dependencies/lib/python2.5/site-packages:\
/opt/pipeline/framework/lib/python2.5/site-packages

engine_ppath = <LOFARSOFTpath>/src/Pulsar/pipeline/recipes/master:\
<LOFARSOFTpath>//src/Pulsar/pipeline/recipes/node: \
<LOFARSOFTpath>/src/Pulsar/pipeline/support:
/opt/pipeline/dependencies/lib/python2.5/site-packages/:\
/opt/pipeline/framework/lib/python2.5/site-packages:\
/opt/LofIm/daily/pyrap/lib:\
/ot/LofIm/daily/lofar/lib/python2.5/site-packages:\
/opt/pythonlibs/lib/python/site-packages

engine_lpath = /opt/pipeline/dependencies/lib:\
/opt/LofIm/daily/pyrap/lib:\
/opt/LofIm/daily/casacore/lib:\
/opt/LofIm/daily/lofar/lib:\
/opt/wcslib/lib/:/opt/hdf5/lib
```

The Pulp specific `tasks.cfg` file is delivered with the Pulp package. The set of "tasks" are named within brackets, their names, the names of available recipes. The `tasks.cfg` file is specified in the `pipeline.cfg` file. The user must edit the `pipeline.cfg` file to reflect the path to this `tasks.cfg` file. Default recipe arguments can be set in the `tasks.cfg` file, seen below. Some recipe defaults appearing in this file are pass-through arguments, and will be recognizable to users as arguments to the respective core components of the "known pulsar pipeline."

```
[buildPulsArch]
recipe     = buildPulsArch
filefactor = 8

[bf2presto]
recipe     = bf2presto
executable = /home/<user>/LOFAR/lofarsoft/release/share/pulsar/bin/bf2presto8
filefactor = 8
collapse   = False
nsigmas    = 7

[buildRSPAll]
```

```
recipe     = buildRSPAll
filefactor = 8

[prepareInf]
recipe     = prepareInf
filefactor = 8

[prepfold]
recipe     = prepfold
executable = /home/<user>/LOFAR/lofarsoft/release/share/pulsar/bin/prepfold
filefactor = 8
nopdsearch = True
nperstokes = 256
noxwin     = True
fine       = True

[rfiplot]
recipe     = rfiplot
executable = /home/<user>/LOFAR/lofarsoft/release/share/pulsar/bin/subdyn.py
filefactor = 8

[bundleFiles]
recipe     = bundleFilesj
filefactor = 8
```

Users will note the explicit paths to `$LOFARSOFT` executables; the configuration parser does not perform shell variable interpolation[20]. Upon Pulp download, users are encouraged to edit these explicit paths to use their own `$LOFARSOFT` build.

2.1.3 Cluster configuraton files

Parallel processing as executed on the LOFAR offline cluster and through the pipeline framework will require three configuration files that define various parameters for job control on the LOFAR offline cluster. None of these files are part of the Pulp distribution. Users will have to locate templates of these files from the generalized pipeline framework location. Once located, these files can be placed arbitrarily, with the arbitrary location specified in the above mentioned and critical `pipeline.cfg` file. Usually, these files are located in the user's specified "runtime_directory", again defined in users' `pipeline.cfg` files.

1. `sub[n].clusterdesc` – Cluster definition file. Tells the framework how to use the cluster.

2. `task.furl` – furl file for tasks to be executed

3. `multiengine furl` – furl file for the multiengine client

Of these files, the user will mostly likely edit the subcluster file, i.e. the `.clusterdesc` file. A "cluster-desc" (cluster description) file is used by the user and the pipeline to assign jobs to the compute cluster, selecting user-specified compute node targets as needed. An example of clusterdesc file content illustrates the configuration options. A typical clusterdesc file content will look like that found in `sub5.clusterdesc`:

```
ClusterName = sub5
# Storage nodes.
Storage.Nodes       = [ lse013..15 ]
Storage.LocalDisks  = [ /data1..4 ]
# Compute nodes.
Compute.Nodes       = [ lce037..45 ]        ==> selected compute nodes
```

---

[20]Troubling.

```
Compute.RemoteDisks  = [ /net/sub1/lse013..15/data1..4 ]
Compute.RemoteFileSys = [ /lse013..15:/data1..4 ]
Compute.LocalDisks    = [ /data ]
# Head nodes.
Head.Nodes            = [ lfe001..2 ]
Head.LocalDisks       = [ /data ]
```

Users can use the range specifier ("..") when altering the "`Compute.Nodes`" parameter, or may explicity enumerate compute nodes in the list, like

```
Compute.Nodes  = [ lce037, lce038, lce45 ], etc..
```

The most common adjustment user's might make here is altering the available "Compute Nodes" list, selecting or deselecting a subset of these "sub5" compute nodes. As specified here, framework job control will (possibly) use all nodes on subnet five.

3 Interfaces

3.1 Pulp command line interfaces

(pulp.py, 1.0, delivered 31.01.2011)

Once a user has established a viable environment for the framework pipeline, execution of Pulp should thence be straight forward.

The Pulp package provides two (2) pre-built command line interfaces ("pre-built" in the sense that users will certainly be writing their own command line interfaces)[21], that is, their own pipeline definitions. The two defintions provided are

1. `pulp.py`

2. `dynspec.py`

The pipeline definition `pulp.py` runs the "basic mode" of the "known pulsar pipeline."[22]. Users should ensure `PATH` access the this program. The "cli" to the executable "pulp.py" module is invoked in the common pythonic idiom,

i.e. `$ python pulp.py -d --job-name ...`[23]

```
 Current usage:

$ pulp.py --d --job-name=<job_name> --obsid=L<yyyy>_<nnnnn[...]>
--arch=<PULP_ARCHIVE> [--pulsar=<pulsar1[,pulsar2,pulsar3]>] [--filefactor=<m>]

where,

--d        = pipeline debug flag, full logging, use.
--obsid    = observation identifier
--job-name = arbitrary job name
--arch     = selected PULSAR ARCHIVE.
--pulsar   = name of pulsar, or csv list of pulsar names.
             eg., --pulsar=B0919+06
                  --pulsar=B0919+06,B0834+06
filefactor = subband splitting factor, <int>
             user specication (range,1--Nsubbands),
             default = 8
```

---

[21] Writing a pipeline definition can be thought of as writing an interface to the Pulp package. Users should be aware that, though the recipes are "independent", there is usually a proxy dependence in that recipes may depend upon the state of the data on disk as it may be provided by prior recipe calls

[22] make_subs_SAS_Ncore_Mmodes.sh in ${LOFARSOFT}/lofarsoft/src/Pulsar/scripts

[23] "Nsubbands" is this example represents the number of subband data files comprising "this" pipeline run.

The second interface provided is `dynspec.py`, and was first describe during the early November (2010) "pulsar busy week". This definition produces only the so-called "rfi report", and a 'png' image of the calculated dynamic spectrum, which is really the product of the pulsar "subdyn.py" script[24]. The Pulp recipe stack is noticeably lighter sans prepfold.

```
Current usage:

$ dynspec.py --d --job-name=<job_name> --obsid=L<yyyy>_<nnnnn[...]>
--arch=<PULP_ARCHIVE> [--pulsar=<pulsar1[,pulsar2,pulsar3]>] [--filefactor=<m>]

where,

--d        = pipeline debug flag, full logging, use.
--obsid    = observation identifier
--job-name = arbitrary job name
--arch     = selected PULSAR ARCHIVE.
--pulsar   = name of pulsar, or csv list of pulsar names.
               eg., --pulsar=B0919+06
                    --pulsar=B0919+06,B0834+06
filefactor = subband splitting factor, <int>
               user specication (range, 1--Nsubbands)
               default = 8
```

The PULP_ARCHIVE Device: Users will find the data products of both `pulp.py` and `dynspec.py` located in the various "RSP[0-248]A" directories defined by the user selected "–arch=archnnn" parameter.[25], in the user selected Pulp archive (`--arch=archnnn`) The user should be aware of the relationship between the defined subclusters of the offline LOFAR cluster, the electable storage nodes, and the location of the data. For example, if a user wishes to process data on subcluster six, i.e. `sub6`, it is incumbent upon the user to specify an appropriate "arch" value, which is primarily determined by data location.
   These values can be found in the PulpEnv() API specification, and are defined thusly:

```
        arch134:  /net/sub5/lse013/data4/PULP_ARCHIVE
        arch144:  /net/sub5/lse014/data4/PULP_ARCHIVE
        arch154:  /net/sub5/lse015/data4/PULP_ARCHIVE
        arch164:  /net/sub6/lse016/data4/PULP_ARCHIVE
        arch174:  /net/sub6/lse017/data4/PULP_ARCHIVE
        arch184:  /net/sub6/lse018/data4/PULP_ARCHIVE
```

As a real world example of a typical "pulp" command line, user will become familiar with this sort of invocation:

```
$ python pulp.py -d --job-name=pulpTest001 --obsid=L2010_09067 --pulsar=B0809+54 \
--arch=144 --filefactor=8 ,
```

where " job-name" will become a directory built in the user's `pipeline_runtime/jobs` directory, and named "job-name". Users will find the framework's "job-name" log file under this jobs directory. This command will push output to the user selected PULP archive on storage node, lse014, and will appear in a path like,

```
/net/sub5/lse014/data4/PULP_ARCHIVE/L2010_09067/
```

Users should expect the final state of the observation directory look roughly like that below.

---

[24]This python script can and should be refactored to be fully funtional on import.
[25]Though Pulp provides much more flexibilty in RSP splitting than the pulsar shell script, users are advised against using large
    numbers for filefactor, especially 248. Mostly because, well, that just seems crazy.

```
/net/sub6/lse018/data4/PULP_ARCHIVE/L2010_09053:
total used in directory 1008988 available 95461848
drwxr-sr-x  [] pulsar       4096 2011-02-02 13:31 .
drwxr-sr-x  [] pulsar         24 2011-02-02 12:58 ..
 -rw-r--r-- [] pulsar 1031558764 2011-02-02 13:31 B1726-00_L2010_09053_plots.tar.gz
 -rw-r--r-- [] pulsar    1502482 2011-02-02 13:31 B1726-00_L2010_09053_pulp.log
drwxr-sr-x  [] pulsar        105 2011-02-02 12:58 incoherentstokes/
 -rw-r--r-- [] pulsar      15018 2011-02-02 12:58 L2010_09053_Master_RSP.list
 -rw-r--r-- [] pulsar      41764 2011-02-02 13:03 L2010_09053.parset
 -rw-r--r-- [] pulsar       1106 2011-02-02 13:03 lofar_default.inf
```

3.2 A note on PulpEnv()and internal package interfaces

The PulpEnv class was designed initially introduced to bring environment variables to compute node recipes. This agency expanded as more processing and observational parameters were introduced in order to simplfy class interfaces within the package. Use of the class helps reduce argument passing at package class interfaces. It should be helping more than it does now. Which just means that the tool is in place, but the interfaces themselves still need some scrubbing.

Though it was implemented *ad hoc* to the initial interface "design", which was essentially – throw everything and anything you need at the call – The PulpEnv class now wraps many parameters, such as those read from observational parset files, and use of the class is implemented.

Here is a map of the current implementation of PulpEnv (support/pulpEnv.py), where observational and processing parameters, and environment variables are made available as a set of instance attributes, defined as strings, unless otherwise indicated.

```
self.obsid        :: LOFAR observation ID, like L<YYYY>_<nnnnn>
self.pulsar       :: Name of targeted pulsar,
self.arch         :: arch, raw archive selection, 'arch134', etc.,
self.subnet       :: subnet in use,
self.environ      :: some user environ values, __dictify(uEnv),
self.LOFARSOFT    :: ${LOFARSOFT},   -- from the head node
self.TEMPO        :: ${TEMPO},       -- from the head node
self.PRESTO       :: ${PRESTO},      -- from the head node
self.oldParsetName:: legacy parset filename 'RTCP.parset.0'
self.parsetPath   :: path to actual parset found.
self.parsetName   :: name of actual parset found.
self.transpose2   :: data through 2nd transpose, <bool>
self.stokes       :: either 'incoherent' or ...
self.archPaths    :: Possible Pulsar Processing Archives
self.pArchive     :: User selected PULP PULSAR ARCHIVE
self.oldLogName   :: old log type: 'run.Storage.R00.log' !! gone.
self.logfilepath  :: full path of log file.
self.obsidPath    :: full path to obsid output.
self.stokesPath   :: full path,  <obsidPath>/[incoherent,raw, ?]/
```

Internal interfaces, however, remain a bit of a mess and could stand some clean up. Had I had the time, I would try to turn a PulpEnv() instance into the only argument passed. Don't know if this is viable, or even desirable, but that was the idea-ish.

<u>Author Note</u>: The implementation of PulpEnv in the pipeline may appear, in a sense, backwards. It should be called at the head node, built there, and passed to the compute nodes. This eliminates multiple reads of the same parset by the compute node jobs. I shall try to exact repair of this awful redundancy. This was part of the necessary growth out of the larval **support/** directory, into the **nodes/** directory and, finally, at long last, to the **master/** directory.

There remain some vestigial attributes in PulpEnv, such as the defunct self.logfilepath. The log files those attributes once addressed have long since disappeared.

3.2.1 A note on Pulp packaging:

Pulp was recently packagized and made importable. Importing Pulp modules now does not require that recipe and support directories appear explicitly in $PYTHONPATH. Packaging is a recent and *ad hoc* supplement that was not available to the Pulp modules at first. Which means that internal imports depend on direct import access to modules. Therefore, it is encouraged that users specify the explicit paths in their $PYTHONPATH because, internally, the package does not have direct access to all Pulp modules without them. These "pipeline" python paths appear as vestigial, a furture development would seek to eliminate this explicit $PYTHONPATH reliance.

Here is an example to further illustrate a Pulp package importation.

```
import pipeline as pulp
```

Packaging provides access to all Pulp recipe and support modules, and can be directly imported. Within python, a user can now access modules in the usual manners, assuming pertinence of the above import.

i.e.,

```
import pulp.support.pulpEnv
```

```
from pipeline import recipes.master.bf2presto
```

```
from pulp.support import pulpEnv as environment
```

Most users will likely not be doing any of this, but is presented here to demonstrate the utility of packaging for module access.

Ultimately, and with minor interface adjustment, such an importable package could be used to launch multiple observation processing jobs.

Eg.,

```
#!/usr/env/python

# Pulp pulsar pipeline controlling script

import pipeline as pulp
from pulp.support import pulpEnv

# no cli, a short ex.
# multiple obsids, filefactors, ...
# dreaming of a day when the user does not
# have to specify unnecessary parameters.

# demonstrate 'filefactor' as 2nd arg.

obsid1     = "L2010_12343"
filefactor1 = 8

obsid2     = "L2010_56748"
filefactor2 = 4

env1 = pulpEnv.PulpEnv(obsid1, filefactor1)
env2 = pulpEnv.PulpEnv(obsid2, filefactor2)

pulp.pulp(env1).go()
pulp.pulp(env2).go()
```

Again, this would require some tweaks to the definition interface, as it is currently implemented. Calling Pulp in such a manner would be equivalent to invoking the pipeline from the command line.

4 Discussion

4.1 Cuisine Exception Handling

One of the more challenging aspects of development in the LOFAR pipeline framework has been the rather unfortunate behaviour of the cuisine library to catch exceptions, raise it's own stripped down "CookError" exception, and toss the traceback of the original source of trouble. This behaviour is ill-advised, but it riddles the WSRTRecipe class and its dependencies. It is advisable that this poor exception "handling" be expunged, and allow actual tracebacks to percolate through the call chain.

4.2 The State of LOFAR parsets

Ideally, LOFAR processing pipelines would draw all observational metadata from the observational parset files. Indeed, the preferred interface to the pipeline would be a single observational argument – the obsid – everything else is or should be available in the parset file. Development of such a model was not feasible, however, as parset files have been under development during this same time. In fact, the reason the pulsar name is a user supplied argument is vestigial, the result of the fact that the "target" parameter was not being populated in the parameter file. It would be nice to get rid of what I consider to be an unnecessary argument. Required and recently-required parameters are now being populated, if somewhat haphazzardly. Bugs have been recently found in parset files, indicating that they are still under development; users and developers should be prepared for unannounced changes in parset content. There is a fair amount of code (in PulpEnv()) devoted to finding moving parsets. With a stable parset location, these methods could be sanitized of this parset hunting.

Location, location, location: A chaotic period of some few weeks saw parset locations migrate on the LOFAR cluster a number of times. At least five (5) paths containing LOFAR parset files were discovered, though the parset location appears to have been stabilized.

4.3 An Investigation of prepfold execution

A report on the investigation into the undocumented behaviour of prepfold within the LOFAR cluster environment. [26]

Prepfold recipe implementation friction was the result of a converegence of unknown behaviour in execution of prepfold itself, with secret files being written to cwd() and then deleted quietly and without notice, and unknown arbitrary file path limits (now know at 100 char). Users and developers should be aware of prepfold's unannounced file writing behaviour.

Here are the temp files produced and removed by `barycenter.c` done during the prepfold process. This effort was undertaken when *barycentering was turned on* during prepfold processing [27]

From `barycenter.c`: [28]

```
remove("tempo.lis");
remove("tempoout_times.tmp");
remove("tempoout_vels.tmp");
remove("resid2.tmp");
remove("bary.tmp");
remove("matrix.tmp");
remove("bary.par");
```

Below is the source of the file read on the TEMPO temp file creation, and specifically, the "`resid2.tmp`" file.

At first, I thought that the embedded (see below) system call to tempo was failing, no $TEMPO in my $PATH. Environment problem on the compute node. fixed. still breaking.

---

[26]Some segments written "in the moment" as notes. Hence the present tense voice at times.

[27](x.inf file: Barycentered? (1=yes, 0=no) = 1

[28]lines 176 - 182 are the best summary of the files written and deleted to be found.)

This is barycenter.c, which writes a temp file for TEMPO called bary.tmp, calls tempo with one argument, the bary.tmp file and pipes the output to another temp file, also called always and everywhere, "tempoout_times.tmp". Currently, these files are being written.

`barycenter.c`, begin line 37:

```
void barycenter (double *topotimes, double *barytimes,
                 double *voverc, long N, char *ra, char *dec, char *obs, char *ephem)
/* This routine uses TEMPO to correct a vector of          */
/* topocentric times (in *topotimes) to barycentric times  */
/* (in *barytimes) assuming an infinite observation         */
/* frequency.  The routine also returns values for the      */
/* radial velocity of the observation site (in units of     */
/* v/c) at the barycentric times.  All three vectors must   */
/* be initialized prior to calling.  The vector length for  */
/* all the vectors is 'N' points.  The RA and DEC (J2000)   */
/* of the observed object are passed as strings in the      */
/* following format: "hh:mm:ss.ssss" for RA and             */
/* "dd:mm:ss.ssss" for DEC.  The observatory site is passed */
/* as a 2 letter ITOA code.  This observatory code must be  */
/* found in obsys.dat (in the TEMPO paths).  The ephemeris  */
/* is either "DE200" or "DE405".                            */

   FILE *outfile;
   long i;
   double fobs = 1000.0, femit, dtmp;
   char command[100], temporaryfile[100];

   /* Write the free format TEMPO file to begin barycentering */

   strcpy(temporaryfile, "bary.tmp");
   outfile = chkfopen(temporaryfile, "w");
   fprintf(outfile, "C  Header Section\n"
           "  HEAD                   \n"
           "  PSR                bary\n"
           "  NPRNT                 2\n"
           "  P0              1.0 1\n"
           "  P1              0.0\n"
           "  CLK          UTC(NIST)\n"
           "  PEPOCH        %19.13f\n"
           "  COORD          J2000\n"
           "  RA              %s\n"
           "  DEC             %s\n"
           "  DM              0.0\n"
           "  EPHEM           %s\n"
           "C  TOA Section (uses ITAO Format)\n"
           "C  First 8 columns must have + or -!\n"
           "  TOA\n", topotimes[0], ra, dec, ephem);

   /* Write the TOAs for infinite frequencies */

   for (i = 0; i < N; i++) {
      fprintf(outfile, "topocen+ %19.13f  0.00     0.0000  0.000000  %s\n",
              topotimes[i], obs);
   }
```

```
   fprintf(outfile, "topocen+ %19.13f  0.00     0.0000  0.000000  %s\n",
           topotimes[N - 1] + 10.0 / SECPERDAY, obs);
   fprintf(outfile, "topocen+ %19.13f  0.00     0.0000  0.000000  %s\n",
           topotimes[N - 1] + 20.0 / SECPERDAY, obs);
   fclose(outfile);

/* Call TEMPO */

/* Check the TEMPO *.tmp and *.lis files for errors when done. */

sprintf(command, "tempo bary.tmp > tempoout_times.tmp");
system(command);

/* Now read the TEMPO results */

strcpy(temporaryfile, "resid2.tmp");

sprintf(command, "tempo bary.tmp > tempoout_times.tmp");   -- FAIL! on no file found.

system(command);
```

As the code presumes, TEMPO produces the file, ''`resid2.tmp`", which prepfold here expects to read. It is at this point that I believe the tempo system call in `barycenter.c` that is not executing properly. Somehow. prepfold makes it impossible to find out what that system call is doing. As mentioned earlier, prepfold runs perfectly well, cutting and pasting the exact command in the shell.

A poorly built system call (is this the nominal PRESTO interface?) that presumes $TEMPO is available directly in your \$PATH, throws away stderr and return signal. According to barycenter.c, tempo ran perfectly, no matter what actually happened.

Another unadvertised feature of `prepfold` and siblings like `barycenter.c` is an arbtrary path limit of 100 characters.

   `barycenter.c`:

```
   temporaryfile[100]
```

Users and other recipe writers should beware long path names when working with elements of PRESTO.

4.4 Open questions

4.4.1 On output paths

In general, production pipelines should maintain tight control over data writing locations. Abiding this prescription, Pulp defines a set of partition specific output archive locations. Malaise best avoided by output control cropped up when user selectable output paths caused the crash of a storage node running multiple pulsar shell script jobs. Such an event ought not happen in a production environment, which is why Pulp defines restricted output archives. Ideally, Pulp would dynamically assign output locations, but that notion seems a tad radical in the current environment.

Nonetheless, the user interface Pulp provides to select an output path is *completely analogous* to how the pulsar pipeline shell script is almost always run.

Writing to local disks.   Initial development requirements for the framework-based pipeline were specified simply as "do what the script does". The nominal operation of the pulsar shell script initially directed, and mostly directs now, output to pre-made Archive directories on designated storage nodes. Actually and originally, it did not do this *per se*, but rather, just wrote stdout directly to the cwd(). In Pulp, these operations were adopted *pro forma*, albeit bent by the obvious need of automated path construction *vis-a-vis* production environment standards.
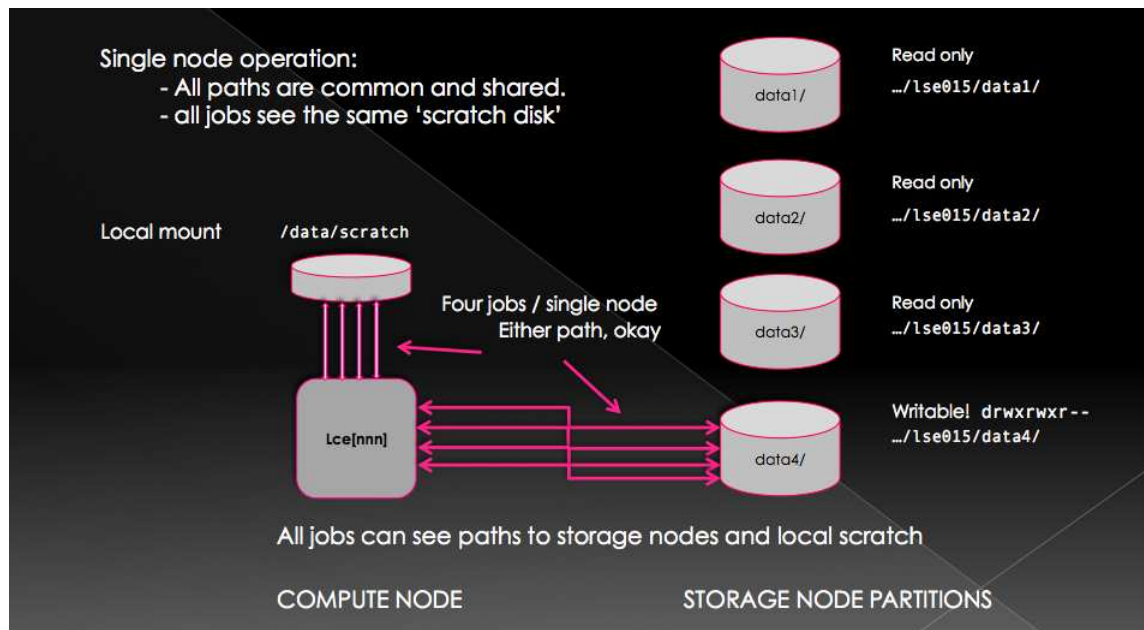
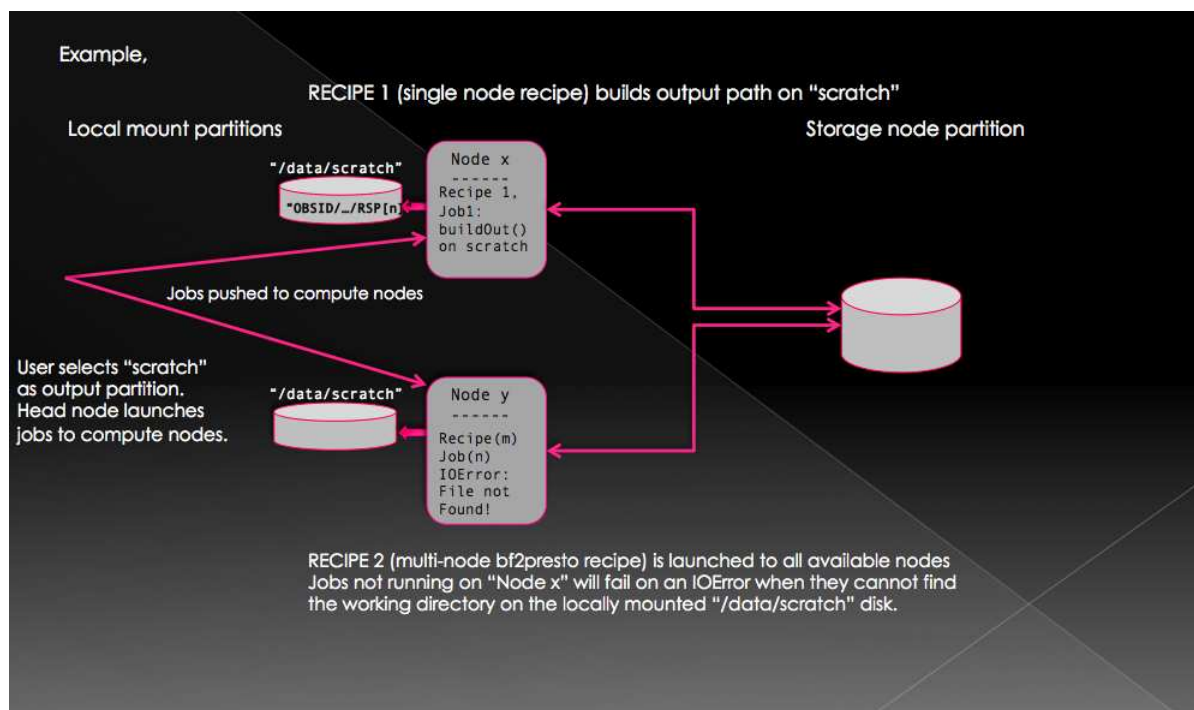Figure 1: Single node operations: all paths are comman and shared

Figure 2: Writing to local disk in a multi-node environment. Paths are common, but not shared.

The *ad hoc* addition of command line options, and specifically that of "–outputpath", was advised against for the Pulp package. The driver to have a full and open output path appears to have been a desire to write to local disk. User's must understand that an output path to local disk would *break* the current Pulp pipeline, unless a user were to restrict pipeline deployment to one (1) compute node.

Figure 1 demonstrates that, for a single node – as the pulsar shell script would operate on – all paths to all write devices are the same for each and every job the pipeline is running. This is certainly *not* true for framework pipeline operations on a multi-node cluster.

The reasons are illustrated in the figure 1. What the figure illustrates can be stated succinctly: when operating in a multi-node cluster compute environment, locally mounted node disks do not share comman pathnames. Writing to local disk would require a complete redesign and would tend toward the programmatic operatons of the standard imaging pipeline.

But even if Pulp were refactored to write to local disk, this still won't work for the pulsar pipeline as it currently operates. Speaking specifically of the "RSPA" processing as it is done now, the current implementation of using symbolic links to "link in" all beam formed data files for "RSPA" will not work. Unless someone knows a way to create symbolic links across or throughout cluster compute nodes, "RSPA" processing cannot be done in a cluster environment. Restricting the whole pipeline to a single node, which is easily doable, would allow "RSPA" processing to go forward, but then that's not exactly cluster computing, is it?

4.5 Future enhancements

5 Acknowledgement

6 Appendix A

6.1 Pulp Classes: a comprehensive listing of public and private methods

The follow images are a schematic representation of Pulp class hierarchy. First, the top level of the package, showing the pipeline definitions, version and package files. (images: `idle` Path browser)[29]
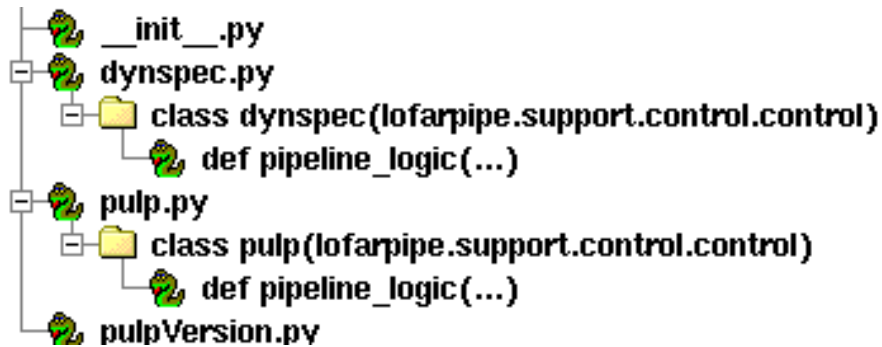

Figure 3: Pulp, v1.0, Class Structure and Methods and Functions, "definition" modules

---

[29]Illustration of an obvious need for a PulpPipeline() base class that would supply the python-private "buildUserEnv()" method will follow shortly. However, once the PulpEnv() call is moved to the master recipes, this method will be deprecated and expunged.

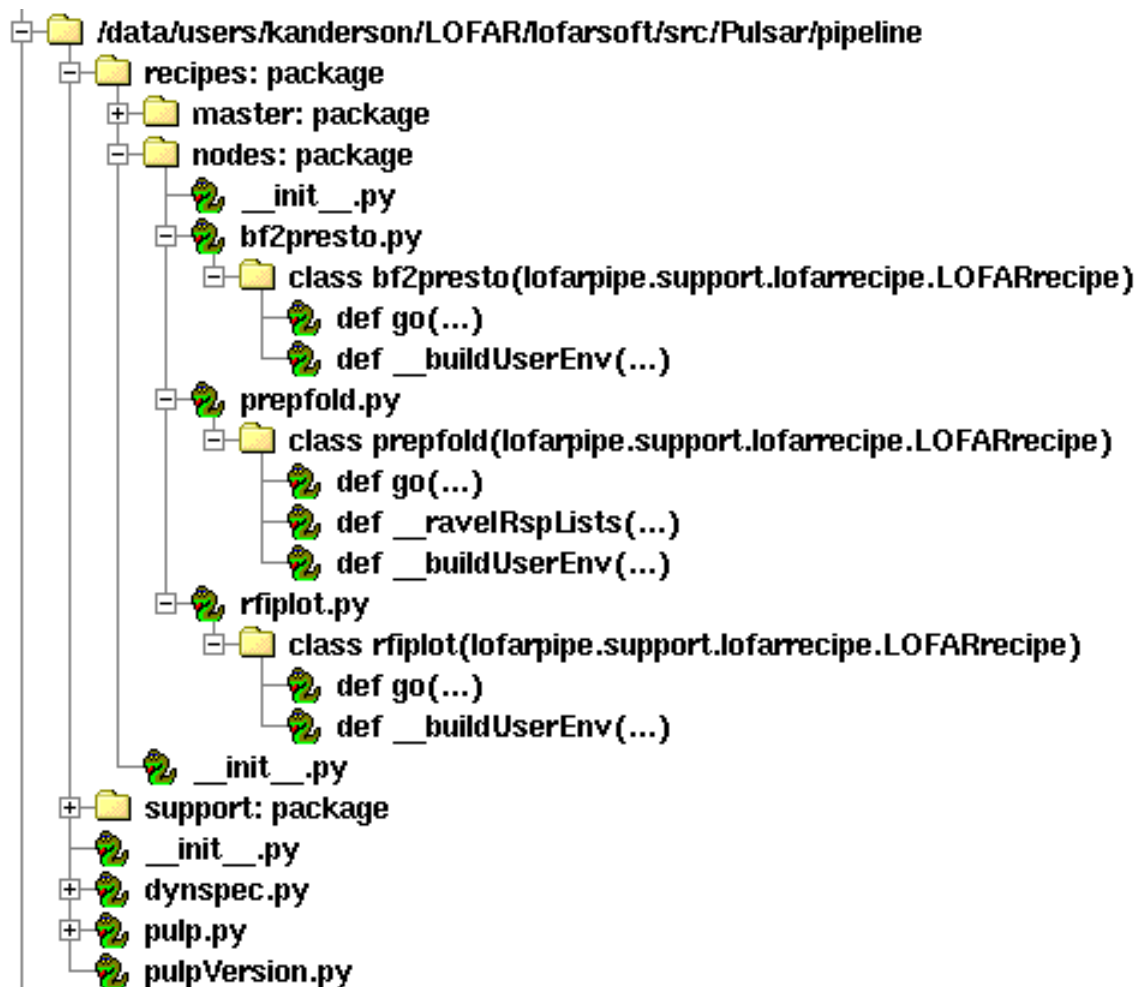Figure 4: Pulp, v1.0, Class Structure, Methods and Functions,'"master" recipe modules

Figure 5: Pulp, v1.0, Hierarchical Class Structure and Methods, "node" modules

Figure 6: Pulp, v1.0, Hierarchical Class Structure and Methods, "support" modules