

# LOFAR User Guide: The Framework Pulsar Pipeline

Pulp, Version 1.0

Document version 1.0

SVN Repository Revision: `--revision--`

K. R. Anderson

SVN Date: 2011-02-02

## Contents

<b>Change record</b>	<b>2</b>
<b>1 Introduction</b>	<b>3</b>
1.1 Applicable Documents . . . . .	3
1.2 Reference Documents . . . . .	4
1.3 Package Overview . . . . .	4
1.3.1 Nominal Pulp Processing . . . . .	5
1.4 Glossary . . . . .	5
<b>2 User Environment</b>	<b>6</b>
2.1 Configuration . . . . .	6
2.1.1 Environment variables, paths . . . . .	6
2.1.2 Framework configuration files . . . . .	7
2.1.3 Cluster configuraton files . . . . .	9
<b>3 Interfaces</b>	<b>10</b>
3.1 The Pulp command line interface . . . . .	10
3.2 A note on PulpEnv() and Pulp package interfaces . . . . .	11
3.2.1 A note on Pulp packaging: . . . . .	12
<b>4 Discussion</b>	<b>13</b>
4.1 Cuisine Exception Handling . . . . .	13
4.2 The State of LOFAR parsets . . . . .	13
4.3 An Investigation of prepfold execution . . . . .	13
4.4 Open questions . . . . .	15
4.4.1 On output paths . . . . .	15
4.5 Future enhancements . . . . .	17
<b>5 Acknowledgement</b>	<b>17</b>

---

<b>6</b>	<b>Appendix A</b>	<b>18</b>
6.1	Pulp Class Hierarchy: a comprehensive listing of public and private methods . . . . .	18

## Change record

ISSUE	DATE	SECTIONS	DESCRIPTION OF CHANGES
1.0	2011-01-31	all	Initial release

# 1 Introduction

>>> import This document is version 1.0 of the LOFAR User Guide for the LOFAR framework pulsar pipeline, known colloquially as ‘pulp’. This User Guide provides general instruction to help users establish a viable user environment as required by the LOFAR pipeline framework in order to run the framework pulsar pipeline in the LOFAR cluster compute environment.

This document will also present details on running the framework pulsar pipeline, known as Pulp (executable: `pulp.py`). Users are encouraged to seek detailed information about the actual processing engaged by the “known pulsar pipeline” from the LOFAR Transient Key Project (TKP) science team. Details regarding LOFAR beam-formed data processing are beyond the scope of this document.<sup>1</sup>

This document assumes some familiarity with terms surrounding the pipeline framework, such as “recipe”, and “head node”, as well as terms associated with “known pulsar pipeline” processing. Such terms will be used throughout this document.

## 1.1 Applicable Documents

Documentation on the Pulp package is provided with the Pulp distribution, and is available with the Pulp package download in `${LOFARSOFT}/src/Pulsar/pipeline/documentation`. Further documentation on the LOFAR pipeline framework, development under the LOFAR pipeline framework, as well as schematic documentation on the Pulp API, are available within the repository itself. Pulp API documents are delivered with the Pulp package as part of any download of the `$LOFARSOFT` repository.

The Pulp current release under the LOFAR USG repository is located in `$LOFARSOFT/src/Pulsar/pipeline`:

```
drwxr-xr-x 15 <user> staff 510 Jan 30 08:21 documentation/
-rw-r--r-- 1 <user> staff 3614 Jan 28 20:33 dynspec.py
-rw-r--r-- 1 <user> staff 1708 Jan 21 17:16 pipeline.cfg
-rwxr-xr-x 1 <user> staff 4545 Jan 28 20:34 pulp.py
-rw-r--r-- 1 <user> staff 40 Jan 28 17:28 pulpVersion.py
drwxr-xr-x 5 <user> staff 170 Nov 1 11:15 recipes/
drwxr-xr-x 22 <user> staff 748 Jan 29 23:09 support/
-rw-r--r-- 1 <user> staff 789 Jan 28 22:07 tasks.cfg
```

In `documentation/`, users will find a directory structure that somewhat mimics the pipeline’s recipe-support directory structure, with recipe API documents found under `./pipeline/documentation/` in ‘master’, ‘nodes’, and ‘support’.

`./Pulsar/pipeline/documentation/:`

-----

```
drwxr-xr-x 15 <user> staff 510 Jan 30 09:11 .
drwxr-xr-x 11 <user> staff 374 Jan 30 08:13 ..
drwxr-xr-x 10 <user> staff 340 Jan 29 15:39 master/
drwxr-xr-x 6 <user> staff 204 Jan 29 15:09 nodes/
drwxr-xr-x 13 <user> staff 442 Jan 29 17:11 support/
```

where a user will find API documents for the pipeline recipes, and all support modules.

`./master:`

-----

```
rw-r--r-- ... bf2presto.html
-rw-r--r-- ... buildPulsArch.html
-rw-r--r-- ... buildRSPAll.html
-rw-r--r-- ... bundleFiles.html
```

<sup>1</sup>Readers of this document who find themselves unfamiliar with the terms, ‘known pulsar pipeline,’ or ‘Jason Hessels’, should probably be reading something else.

```

-rw-r--r-- ... prepareInf.html
-rw-r--r-- ... prepfold.html
-rw-r--r-- ... rfipplot.html

./nodes:
-----
-rw-r--r-- ... bf2presto.html
-rw-r--r-- ... prepfold.html
-rw-r--r-- ... rfipplot.html

./support:
-----
  rw-r--r-- ... RSPlist.html
-rw-r--r-- ... bf2Pars.html
-rw-r--r-- ... buildRSPS.html
-rw-r--r-- ... bundlePlots.html
-rw-r--r-- ... foldingData.html
-rw-r--r-- ... fullRSP.html
-rw-r--r-- ... pardata.html
-rw-r--r-- ... prepInfFiles.html
-rw-r--r-- ... pulpEnv.html
-rw-r--r-- ... rfiDirectories.html

```

These are accessible and viewable through any browser.

## 1.2 Reference Documents

For documentation on the framework itself, as well as guidance on recipe writing, the user is invited to examine the framework documentation available in the repository under

`${LOFARSOFT}/src/pipeline/docs` ,

A subset of the framework’s API is documented in the Pulp API documentation described above.

## 1.3 Package Overview

The Pulp package is based upon the the LOFAR TKP Pulsar Group’s pipelining shell script<sup>2</sup>. Pulp is built on the LOFAR pipeline framework and, as with all such pipelines, comprises a set of ‘recipes’ that are executed in successive order. In a cluster compute environment, these recipes can be paired in order to facilitate parallel processing on a cluster and/or subsections of a cluster. Depending upon the required action, recipes may or may not be paired across head node/compute node connections.<sup>3</sup> When a “head node” recipe is not paired with a matching “compute node” recipe (name matching of head/compute node recipes is required by the framework), this necessarily implies that a single process will be activated on a compute node in order to perform a task that does not require parallelization. For example, building the storage node directory structure for an observation.

In order to illustrate this meaning, Table 1 presents the overall organisation of the Pulp pulsar pipeline package. The table demonstrates those recipes that will perform the tasks requiring multiple parallel jobs. Those recipes (bolded) show matching named “node recipe” modules in the package. It is these head node/compute node recipe pairs that multiply execute all requested processing jobs. Those jobs, and the resulting job queues, are arranged and built by the respective head node recipe, and then farmed to the cluster based upon a user’s “clusterdesc” file (see §2).

<sup>2</sup>`make_subs_SAS_Ncore_Mmodes.sh` in `${LOFARSOFT}/lofarsoft/src/Pulsar/scripts`

<sup>3</sup>This is explicitly true for the Pulp package, as it employs the framework’s IPython facility. This statement is likely not correct for the “new style” imaging pipeline, which employs the more recently introduced ssh method.

MASTER RECIPE	NODE RECIPE	SUPPORT MODULES
buildPulsArch	—	buildRSPS, RSPlist
<b>bf2presto</b>	<b>bf2presto</b>	bf2Pars, pulpEnv
buildRSPALL	—	fullRSP, pulpEnv
prepareInf	—	prepareInfFiles, pulpEnv
<b>prepfold</b>	<b>prepfold</b>	pulpEnv
<b>rfiplot</b>	<b>rfiplot</b>	rfiDirectories, pulpEnv
bundleFiles	—	bundlePlots, pulpEnv

Table 1: Pulp distribution and recipe relations

Note: **support/** modules may support other **support/** modules. Recipes do not support other recipes. As indicated earlier (§1), the described Pulp package is delivered via checkout of the LOFAR USG code repository.

### 1.3.1 Nominal Pulp Processing

The nominal processing performed by the Pulp package involves,

- The number of subbands in a observation is entirely selected by the observer, however, nominally 248 subbands of ‘incoherent’ beam-formed data are usually delivered and processed. Currently, Pulp can only process so-called ‘incoherentstokes’ beam-formed data. It was mandated that development on ‘coherentstokes’, or possibly ‘stokes’, perhaps even ‘raw’, data processing would be tabled until coherent stokes data are being written to the cluster in the new LOFAR HDF5 format<sup>4</sup>
- The pipeline is capable of handling any user selected “splitting” factor, i.e. the ‘filefactor’ parameter within Pulp. This factor is known as “ncores” in the pulsar shell script. Pulp places no restriction on what ‘filefactor’ can be, though nominal processing will usually dictate the default `filefactor = 8`.
- In the special case that a user select `filefactor == 1`, No “all” processing will be done. That is, no RSPA is made, and no RSPA processing, *per se*, occurs. Output will appear in the expected ‘RSP0’ directory. RSP0 is equivalent to RSPA, when `filefactor=1`.
- The shell script “-rfi” switch is not a switch with Pulp (it could be made that way). Pulp executes the relevant steps automatically and the product, a “.rffireport” file, is bundled as a standard pipeline data product.

Details on the Pulp command line interface are provided in §3, “Interfaces”.

## 1.4 Glossary

- **API** – Application programming interface (!Anton Pannekoek Instituut).
- **Compute node** – One of an assigned number of cluster nodes configured for computing. For LOFAR, these compute nodes have NFS connections to selected storage nodes, which is where both input and output data will be read from and written to. Within the context of the LOFAR pipeline framework, compute nodes are often referred to simply as “nodes”.
- **Cluster** – A multi-node computing cluster.
- **Head node** – The head (controlling) node of a cluster environment. In the LOFAR computing environment, the head node is lfe001. (rumours of a second “head node” on the LOFAR cluster, “lfe002”, have been reported, though remain confirmed by this user)
- **Pulp** – the “known pulsar pipeline” implemented under the LOFAR pipeline framework.
- **TKP**– the Transient Key Project (LOFAR).

<sup>4</sup>see LOFAR document, *LOFAR-USG-ICD-003, Beam-formed Data*, Alexov et al, ASTRON, 2010.

## 2 User Environment

In computing environments, pipelines function as automated, simulated users at the command line. More aptly, a pipeline functions as a command line user agent. Therefore, and in order that the Pulp package execute properly, Pulp users will need to take a number of steps in adjusting their compute environments. Framework pipeline operations require that a number of locally defined configuration files be accessible to the framework. These files are (in no particular order):

1. `pipeline.cfg` – configuration for framework operations
2. `tasks.cfg` – define things to do
3. `sub[n].clusterdesc` – define target compute nodes
4. `task.furl` – job control
5. `multiengine.furl` – ip engine control

The two configuration files, `pipeline.cfg` and `tasks.cfg`, listed here are provided with the Pulp download. Described below, users will want to edit the pipeline configuration file to use their own builds of the \$LOFARSOFT code repository. Users may edit the `tasks.cfg` file, which is where a user may easily change default arguments of any or all recipes. In all likelihood, most users will not want or need to do this. These configuration files, and others, are discussed in detail in §2.1.2, “Framework configuration files”.

### 2.1 Configuration

#### 2.1.1 Environment variables, paths

The following steps should be taken by users to establish the environment variables and directories required by the pipeline framework.

- Define LOFARSOFT, LOFARROOT (`=/opt/LofIm/daily/lofar`)
- `source ${LOFARSOFT}/devel_common/scripts/init.sh`
- `Use LofIm`
- Make a “`pipeline_runtime/`” directory (arbitrary location)
- Edit user’s `pipeline.cfg` file to reflect this path
- Adjust \$PYTHONPATH to include framework python libraries.

Invocation of ‘`Use LofIm`’ will define the environment variables, \$TEMPO, and \$PRESTO, needed for embedded PRESTO<sup>5</sup> operations.

Users will need to make a directory the framework expects to use for pipeline activities, such as writing various log files. Within the above mentioned configuration file, ‘`pipeline.cfg`’, discussed in detail below, users will define the configuration parameter, ‘`runtime_directory`’.

i.e.

```
runtime_directory = /path/to/your/runtime_directory
```

As illustrated, this directory is arbitrary in both name and location, though by convention, it is usually located under a user’s home directory.<sup>6</sup>

Once these are defined (usually within a user’s shell resource file), a user’s \$PYTHONPATH should further include the following paths described below (defer to J. Swinbank if any of this changes or is unneeded.)

<sup>5</sup><http://www.cv.nrao.edu/~sransom/presto>

<sup>6</sup>Explicit paths must appear in this configuration file.

```

PYTHONPATH=/opt/pipeline/dependencies/lib/python2.5/site-packages:\
/opt/pipeline/framework/lib/python2.5/site-packages:          \
/opt/LofIm/daily/pyrap/lib:                                   \
${LOFARROOT}/lib/python2.5/site-packages:                     \
/opt/pythonlibs/lib/python/site-packages:                     \
${LOFARSOFT}/src/Pulsar/pipeline/recipes/master:              \
${LOFARSOFT}/src/Pulsar/pipeline/recipes/nodes:               \
${LOFARSOFT}/src/Pulsar/pipeline/support

```

A user's \$PATH environment variable should be adjusted to include the following paths.

\$PATH include:

```

${LOFARSOFT}/release/bin:\
${LOFARSOFT}/release/share/pulsar/bin:\
${LOFARSOFT}/src/Pulsar/pipeline:\
/opt/LofIm/daily/casarest/bin:\
/opt/pipeline/dependencies/bin:\
/opt/LofIm/daily/askapsoft/bin:\
/opt/scripts:${PATH}

```

N.B. These paths are subject to future alteration of the LOFAR system.

### 2.1.2 Framework configuration files

The pipeline framework functions through the use and availability of a limited set of configuration files, i.e., .cfg files. These configuration files are listed and briefly described. For more information, see the documentation on the pipeline framework (Swinbank, 2011).

pipeline.cfg  
tasks.cfg

A user will notice that these two configuration files are delivered by the svn repository. These files are required by the framework. They are *not* part of the Pulp package. Though they must be tailored by the user to their particular environment and file system, the files are provided by the repository as both convenience and necessity; the `pipeline.cfg` the convenience, the `tasks.cfg` the necessity.

The `pipeline.cfg` file is a “pipeline generic” configuration file applicable to all framework pipelines. The Pulp `tasks.cfg` file is not. It is advised that this file not be altered without perspicacious intent.

Here is the top level of a delivered \$LOFARSOFT Pulp directory:

```

[...]/lofarsoft/src/Pulsar/pipeline:
drwxr-xr-x      .
drwxr-xr-x      ..
drwxr-xr-x      .svn
-rw-r--r--@    ... __init__.py
drwxr-xr-x      ... documentation/
-rw-r--r--      ... dynspec.py
-rw-r--r--      ... pipeline.cfg  ---> user paths and locations
-rwxr-xr-x      ... pulp.py
-rw-r--r--      ... pulpVersion.py
drwxr-xr-x      ... recipes/
drwxr-xr-x      ... support/
-rw-r--r--      ... tasks.cfg    ---> task/recipe definitions

```

The reasons these particular files are delivered as part of the Pulp package are two fold: provide user's with a configuration basis, and to specify Pulp functionality, delivered by the `tasks.cfg` file. This configuration file is critical to proper pipeline operations and should not be altered by users unless done so with understanding of pipeline operations.

Here is a quasi-generic pipeline configuration file:



```

[DEFAULT]
runtime_directory = /some/path/pipeline_runtime
recipe_directories = [<LOFARSOFTpath>/src/Pulsar/pipeline/recipes]
lofarroot          = /opt/LofIm/daily/lofar
default_working_directory = /data/scratch/<user>
task_files          = [<LOFARSOFTpath>/src/Pulsar/pipeline/tasks.cfg]

[layout]
job_directory      = %(runtime_directory)s/jobs/%(job_name)s
log_directory      = %(job_directory)s/logs/
vds_directory      = %(job_directory)s/vds
parset_directory   = %(job_directory)s/parsets
results_directory  = %(job_directory)s/results/%(start_time)s

[cluster]
clustername        = pulsar
#clusterdesc       = %(runtime_directory)s/sub5.clusterdesc
clusterdesc        = %(runtime_directory)s/line_runtime/sub6.clusterdesc
task_furl          = %(runtime_directory)s/task.furl
multiengine_furl   = %(runtime_directory)s/multiengine.furl

[deploy]
script_path        = /opt/pipeline/framework/bin
controller_ppath    = <LOFARSOFTpath>/src/Pulsar/pipeline/support:\
/opt/pipeline/dependencies/lib/python2.5/site-packages:\
/opt/pipeline/framework/lib/python2.5/site-packages

engine_ppath = <LOFARSOFTpath>/src/Pulsar/pipeline/recipes/master:\
<LOFARSOFTpath>/src/Pulsar/pipeline/recipes/node: \
<LOFARSOFTpath>/src/Pulsar/pipeline/support:
/opt/pipeline/dependencies/lib/python2.5/site-packages:\
/opt/pipeline/framework/lib/python2.5/site-packages:\
/opt/LofIm/daily/pyrap/lib:\
/opt/LofIm/daily/lofar/lib/python2.5/site-packages:\
/opt/pythonlibs/lib/python/site-packages

engine_lpath = /opt/pipeline/dependencies/lib:\
/opt/LofIm/daily/pyrap/lib:\
/opt/LofIm/daily/casacore/lib:\
/opt/LofIm/daily/lofar/lib:\
/opt/wcslib/lib:/opt/hdf5/lib

```

The Pulp specific `tasks.cfg` file is delivered with the Pulp package. The set of “tasks” are named within brackets, their names, the names of available recipes. The `tasks.cfg` file is specified in the `pipeline.cfg` file. Default recipe argument values can be set in the `tasks.cfg` file, as can be seen below.

```

[buildPulsArch]
recipe      = buildPulsArch
filefactor  = 8

```

```

[bf2presto]

```

```

recipe      = bf2presto
executable  = /home/kanderson/LOFAR/lofarsoft/release/share/pulsar/bin/bf2presto8
filefactor  = 8
collapse    = False
nsigmas     = 7

[buildRSPAll]
recipe      = buildRSPAll
filefactor  = 8

[prepareInf]
recipe      = prepareInf
filefactor  = 8

[prepfold]
recipe      = prepfold
executable  = /home/kanderson/LOFAR/lofarsoft/release/share/pulsar/bin/prepfold
filefactor  = 8
nopdsearch  = True
nperstokes  = 256
noxwin      = True
fine        = True

[rfiplot]
recipe      = rfiplot
executable  = /home/kanderson/LOFAR/lofarsoft/release/share/pulsar/bin/subdyn.py
filefactor  = 8

[bundleFiles]
recipe      = bundleFilesj
filefactor  = 8

```

Users will note the explicit paths to `$LOFARSOFT` executables. This is because the configuration parser does not perform shell variable interpolation. Upon Pulp download, users are encouraged to edit these explicit paths to use their own `$LOFARSOFT` build.

### 2.1.3 Cluster configuraton files

Parallel processing as executed on the LOFAR offline cluster and through the pipeline framework will require three configuration files that define various parameters for job control on the LOFAR offline cluster. None of these files are part of the Pulp distribution. Users will have to locate templates of these files from the generalized pipeline framework location. Once located, these files can be placed arbitrarily, with the arbitrary location specified in the above mentioned and critical `pipeline.cfg` file. Usually, these files are located in the user's specified "runtime\_directory", again defined in users' `pipeline.cfg` files.

1. `sub[n].clusterdesc` – Cluster definition file. Tells the framework how to use the cluster.
2. `task.furl` – furl file for tasks to be executed
3. `multiengine furl` – furl file for the multiengine client

Of these files, the user will mostly likely edit the subcluster file, selecting compute node targets as needed. An example of clusterdesc file content illustrates the configuration options.

```
sub5.clusterdesc:
```

```
ClusterName = sub5
```

```
# Storage nodes.
Storage.Nodes      = [ lse013..15 ]
Storage.LocalDisks = [ /data1..4 ]
# Compute nodes.
Compute.Nodes      = [ lce037..45 ]      ==> selected compute nodes
Compute.RemoteDisks = [ /net/sub1/lse013..15/data1..4 ]
Compute.RemoteFileSys = [ /lse013..15:/data1..4 ]
Compute.LocalDisks  = [ /data ]
# Head nodes.
Head.Nodes         = [ lfe001..2 ]
Head.LocalDisks     = [ /data ]
```

The most common adjustment user's might make here is altering the available "Compute Nodes" list, selecting or deslecting a subset of these "sub5" compute nodes. As specified here, framework job control will (possibly) use all nodes on subnet five.

## 3 Interfaces

### 3.1 The Pulp command line interface

(pulp.py, 1.0, delivered 31.01.2011)

Once a user has established a viable environment for the framework pipeline, execution of Pulp should thence be straight forward.

The command line interface to the executable "pulp.py" module is described below. If the user does not have their \$PATH environment set to access pulp.py, then the command line below should be invoked in the common pythonic idiom,

i.e. \$ python pulp.py -d --job-name ...

Current usage:

```
$ pulp.py --d --job-name=<job_name> --arch=<PULP_ARCHIVE>
[--pulsar=<pulsar1[,pulsar2,pulsar3,]>] [--filefactor=<m>]
```

where,

```
--d          = pipeline debug flag, full logging, use.
--obsid      = observation identifier
--job-name   = arbitrary job name
--arch       = selected PULSAR ARCHIVE.
--pulsar     = name of pulsar, or csv list of pulsar names.
               eg., --pulsar=B0919+06
                   --pulsar=B0919+06,B0834+06
                   (default=FOV pulsar name search)
filefactor   = subband splitting factor, <int>
               optional user specification (range, 1-248)
               default = 8
```

The user should be aware of the relationship between the defined subclusters on the offline LOFAR cluster, the selectable storage nodes, and the location of the data. For example, if a user wishes to process data on the subcluster six, i.e. sub6, it is incumbent upon the user to specify an appropriate "arch" value, which is primarily determined by data location.

These values can be found in the PulpEnv() API specification, and defined thusly:

```
arch134:  /net/sub5/lse013/data4/PULP_ARCHIVE
arch144:  /net/sub5/lse014/data4/PULP_ARCHIVE
```

```

arch154: /net/sub5/lse015/data4/PULP_ARCHIVE
arch164: /net/sub6/lse016/data4/PULP_ARCHIVE
arch174: /net/sub6/lse017/data4/PULP_ARCHIVE
arch184: /net/sub6/lse018/data4/PULP_ARCHIVE

```

As a real world example of a typical “pulp” command line, user will become use to this sort of invocation.

```
$ python pulp.py -d --job-name=pulpTest001 --obsid=L2010_09067 --pulsar=B0809+54 \
--arch=144 --filefactor=8 ,
```

where “job-name” will become a directory built in the user’s `pipeline_runtime/jobs` directory, and named “job-name”. Users will find the framework’s “job-name” log file under this jobs directory. This command will push output to the user selected PULP archive on storage node, lse014, and will appear in a path like,

```
/net/sub5/lse014/data4/PULP_ARCHIVE/L2010_09067/
```

Where the final state of the observation directory will look roughly like this.

```

/net/sub6/lse018/data4/PULP_ARCHIVE/L2010_09053:
total used in directory 1008988 available 95461848
drwxr-sr-x pulsar      4096 2011-02-02 13:31 .
drwxr-sr-x pulsar      24 2011-02-02 12:58 ..
-rw-r--r-- pulsar 1031558764 2011-02-02 13:31 B1726-00_L2010_09053_plots.tar.gz
-rw-r--r-- pulsar  1502482 2011-02-02 13:31 B1726-00_L2010_09053_pulp.log
drwxr-sr-x pulsar      105 2011-02-02 12:58 incoherentstokes/
-rw-r--r-- pulsar  15018 2011-02-02 12:58 L2010_09053_Master_RSP.list
-rw-r--r-- pulsar  41764 2011-02-02 13:03 L2010_09053.parset
-rw-r--r-- pulsar   1106 2011-02-02 13:03 lofar_default.inf

```

### 3.2 A note on PulpEnv() and Pulp package interfaces

The PulpEnv class was designed initially introduced to bring environment variables to compute node recipes. This agency expanded as more processing and observational parameters were introduced in order to simplify class interfaces within the package. Use of the class helps reduce argument passing at package class interfaces. It should be helping more than it does now. Which just means that the tool is in place, but the interfaces themselves still need some scrubbing.

Though it was implemented *ad hoc* to the initial interface “design”, which was essentially – throw everything and anything you need at the call – The PulpEnv class now wraps many parameters, such as those read from observational parset files, and use of the class is implemented.

Here is a map of the current implementation of PulpEnv (support/pulpEnv.py), where observational and processing parameters, and environment variables are made available as a set of instance attributes, defined as strings, unless otherwise indicated.

```

self.obsid      :: LOFAR observation ID, like L<YYYY>_<nnnnn>
self.pulsar     :: Name of targeted pulsar,
self.arch       :: arch, raw archive selection, 'arch134', etc.,
self.subnet     :: subnet in use,
self.envIRON    :: some user environ values, __dictify(uEnv),
self.LOFARSOFT  :: ${LOFARSOFT}, -- from the head node
self.TEMPO      :: ${TEMPO}, -- from the head node
self.PRESTO     :: ${PRESTO}, -- from the head node
self.oldParsetName:: legacy parset filename 'RTCP.parset.0'
self.parsetPath  :: path to actual parset found.
self.parsetName  :: name of actual parset found.
self.transpose2  :: data through 2nd transpose, <bool>
self.stokes     :: either 'incoherent' or ...

```

```

self.archPaths      :: Possible Pulsar Processing Archives
self.pArchive       :: User selected PULP PULSAR ARCHIVE
self.oldLogName      :: old log type: 'run.Storage.R00.log' !! gone.
self.logfilepath     :: full path of log file.
self.obsidPath       :: full path to obsid output.
self.stokesPath      :: full path, <obsidPath>/[incoherent,raw, ?]/

```

Internal interfaces, however, remain a bit of a mess and could stand some clean up. Had I had the time, I would try to turn a PulpEnv() instance into the only argument passed. Don't know if this is viable, or even desirable, but that was the idea-ish.

*Author Note:* The implementation of PulpEnv in the pipeline may appear, in a sense, backwards. It should be called at the head node, built there, and passed to the compute nodes. This eliminates multiple reads of the same parset by the compute node jobs. I shall try to exact repair of this awful redundancy. This was part of the necessary growth out of the larval **support/** directory, into the **nodes/** directory and, finally, at long last, to the **master/** directory.

There remain some vestigial attributes in PulpEnv, such as the defunct self.logfilepath. The log files those attributes once addressed have long since disappeared.

### 3.2.1 A note on Pulp packaging:

The Pulp package has been recently made an importable package, and importing Pulp now does not demand that recipe and support directories appear explicitly in \$PYTHONPATH. Since packaging is a recent and ad hoc supplement that was not available to the Pulp modules at first. Which means that internal imports depend on direct import access to modules. Therefore, it is encouraged that users specify the explicit paths in their \$PYTHONPATH because, internally, the package does not have direct access to all Pulp modules without them. These “pipeline” python paths appear as vestigial, a future development would seek to eliminate this explicit \$PYTHONPATH reliance.

Here is an example to further illustrate a Pulp package importation.

```
import pipeline as pulp
```

Packaging provides access to all Pulp recipe and support modules, and can be directly imported. Within python, a user can now access modules in the usual manners, assuming pertinence of the above import. i.e.,

```
import pulp.support.pulpEnv
```

```
from pipeline import recipes.master.bf2presto
```

```
from pulp.support import pulpEnv as environment
```

Most users will likely not be doing any of this, but is presented here to demonstrate the utility of packaging for module access.

Ultimately, and with minor interface adjustment, such an importable package could be used to launch multiple observation processing jobs.

Eg.,

```
#!/usr/env/python
```

```
import pipeline as pulp
```

```
# multiple obsid and pulsar values, different filefactors, ...
```

```
obsid1      = "L2010_12343"
```

```
pulsar1     = "B0809+54"
```

```
filefactor1 = 8
```

```

obsid2      = "L2010_56748"
pulsar2     = "B0111-11"
filefactor2 = 4

pulp.pulp(obsid1,pulsar1,filefactor1).go()
pulp.pulp(obsid2,pulsar2,filefactor2).go()

```

Again, this would require some tweaks to the definition interface, as it is currently implemented. Calling Pulp in such a manner would be equivalent to invoking the pipeline from the command line.

## 4 Discussion

### 4.1 Cuisine Exception Handling

One of the more challenging aspects of development in the LOFAR pipeline framework has been the rather unfortunate behaviour of the cuisine library to catch exceptions, raise it's own stripped down "CookError" exception, and toss the traceback of the original source of trouble. This behaviour is ill-advised, but it riddles the WSRTRecipe class and its dependencies. It is advisable that this poor exception "handling" be expunged, and allow actual tracebacks to percolate through the call chain.

### 4.2 The State of LOFAR parsets

Ideally, LOFAR processing pipelines would draw all observational metadata from the observational parset files. Indeed, the preferred interface to the pipeline would be a single observational argument – the obsid – everything else is or should be available in the parset file. Development of such a model was not feasible, however, as parset files have been under development during this same time. In fact, the reason the pulsar name is a user supplied argument is vestigial, the result of the fact that the "target" parameter was not being populated in the parameter file. It would be nice to get rid of what I consider to be an unnecessary argument. Required and recently-required parameters are now being populated, if somewhat haphazardly. Bugs have been recently found in parset files, indicating that they are still under development; users and developers should be prepared for unannounced changes in parset content. There is a fair amount of code (in PulpEnv()) devoted to finding moving parsets. With a stable parset location, these methods could be sanitized of this parset hunting.

Location, location, location: A chaotic period of some few weeks saw parset locations migrate on the LOFAR cluster a number of times. At least five (5) paths containing LOFAR parset files were discovered, though the parset location appears to have been stabilized.

### 4.3 An Investigation of prepfold execution

A report on the investigation into the undocumented behaviour of prepfold within the LOFAR cluster environment, various other webs of intrigue. Prefold recipe implementation friction was the result of a convergence of undocumented behaviours by both the Framework and the execution of prepfold itself, with secret files being written to cwd() and then deleted quietly and without notice. That the framework did not pass the environment to compute nodes is not unexpected, just not documented – at least not then. Ultimately, the environment is not a framework problem per se, but users and developers should be aware of prepfold's secret file writing behaviour.

Firstly, here are the temp files produced and removed by `barycenter.c` done during the prepfold process. This effort was undertaken when *barycentering was turned on* during prepfold processing (x.inf file: barycentering? 1)

`barycenter.c`, lines 176 - 182 (the best summary of this secret behaviour to be found.)

```

remove("tempo.lis");
remove("tempoout_times.tmp");

```

```

remove("tempoout_vels.tmp");
remove("resid2.tmp");
remove("bary.tmp");
remove("matrix.tmp");
remove("bary.par");

```

Below is the source of the file read on the TEMPO temp file creation, and specifically, the "resid2.tmp" file.

At first, I thought that the embedded (see below) system call to tempo was failing, that I did not have tempo in my \$PATH. But that can't be right. Because I am able to run prepfold perfectly well directly in the shell on the nodes, to completion, no problems.

Except, the prepfold behaviour looks exactly like the tempo call is failing in some way.

This is barycenter.c, which writes a temp file for TEMPO called bary.tmp, calls tempo with one argument, the bary.tmp file and pipes the output to another temp file, also called always and everywhere, "tempoout\_times.tmp". Currently, these files are being written.

barycenter.c, begin line 37:

```

void barycenter (double *topotimes, double *barytimes,
                 double *voverc, long N, char *ra, char *dec, char *obs, char *ephem)
/* This routine uses TEMPO to correct a vector of          */
/* topocentric times (in *topotimes) to barycentric times  */
/* (in *barytimes) assuming an infinite observation       */
/* frequency. The routine also returns values for the     */
/* radial velocity of the observation site (in units of   */
/* v/c) at the barycentric times. All three vectors must */
/* be initialized prior to calling. The vector length for */
/* all the vectors is 'N' points. The RA and DEC (J2000) */
/* of the observed object are passed as strings in the    */
/* following format: "hh:mm:ss.ssss" for RA and          */
/* "dd:mm:ss.ssss" for DEC. The observatory site is passed */
/* as a 2 letter ITOA code. This observatory code must be */
/* found in obsys.dat (in the TEMPO paths). The ephemeris */
/* is either "DE200" or "DE405".                          */

FILE *outfile;
long i;
double fobs = 1000.0, femit, dtmp;
char command[100], temporaryfile[100];

/* Write the free format TEMPO file to begin barycentering */

strcpy(temporaryfile, "bary.tmp");
outfile = chkfopen(temporaryfile, "w");
fprintf(outfile, "C Header Section\n"
               "  HEAD                \n"
               "  PSR                  bary\n"
               "  NPRNT                2\n"
               "  P0                   1.0 1\n"
               "  P1                   0.0\n"
               "  CLK                  UTC(NIST)\n"
               "  PEPOCH               %19.13f\n"
               "  COORD                J2000\n"
               "  RA                   %s\n"
               "  DEC                  %s\n"

```

```

"   DM                                0.0\n"
"  EPHEM                             %s\n"
"C   TOA Section (uses ITAO Format)\n"
"C   First 8 columns must have + or -!\n"
"   TOA\n", topotimes[0], ra, dec, ephemeris);

/* Write the TOAs for infinite frequencies */

for (i = 0; i < N; i++) {
    fprintf(outfile, "topocen+ %19.13f  0.00      0.0000  0.000000  %s\n",
              topotimes[i], obs);
}
fprintf(outfile, "topocen+ %19.13f  0.00      0.0000  0.000000  %s\n",
          topotimes[N - 1] + 10.0 / SECPERDAY, obs);
fprintf(outfile, "topocen+ %19.13f  0.00      0.0000  0.000000  %s\n",
          topotimes[N - 1] + 20.0 / SECPERDAY, obs);
fclose(outfile);

/* Call TEMPO */

/* Check the TEMPO *.tmp and *.lis files for errors when done. */

sprintf(command, "tempo bary.tmp > tempoout_times.tmp");
system(command);

/* Now read the TEMPO results */

strcpy(temporaryfile, "resid2.tmp");

sprintf(command, "tempo bary.tmp > tempoout_times.tmp");  -- FAIL! on no file found.

system(command);

```

As the code presumes, TEMPO produces the file, ‘‘resid2.tmp”, which prepfold here expects to read. It is at this point that I believe the tempo system call in `barycenter.c` that is not executing properly. Somehow, prepfold makes it impossible to find out what that system call is doing. As mentioned earlier, prepfold runs perfectly well, cutting and pasting the exact command in the shell.

A poorly built system call (! is this the way all of presto interacts?) that presumes \$TEMPO is available directly in your %\$PATH, throws away stderr and any return signal. According to `barycenter.c`, tempo ran perfectly, no matter what actually happened.

Another unadvertised feature of `prepfold` and siblings like `barycenter.c` is an arbitrary filename size limited to 100 characters.

`barycenter.c`:

```
temporaryfile[100]
```

Users and other recipe writers should beware long path names when playing with elements of PRESTO.

## 4.4 Open questions

### 4.4.1 On output paths

In general, production pipelines should maintain tight control over data writing locations. Abiding this prescription, Pulp defines a set of partition specific output archive locations. This seems to be a strange notion at ASTRON, but malaise best avoided by output control cropped up when user selectable output



paths caused the crash of a storage node running multiple pulsar shell script jobs. Such an event ought not happen in a production environment, which is why Pulp defines restricted output archives. Ideally, Pulp would dynamically assign output locations, but that notion seems a tad radical in the current environment.

Nonetheless, the user interface Pulp provides to select an output path is *completely analogous* to how the pulsar pipeline shell script is almost always run.

**Writing to local disks.** Initial development requirements for the framework-based pipeline were specified simply as “do what the script does”. The nominal operation of the pulsar shell script initially directed, and mostly directs now, output to pre-made Archive directories on designated storage nodes. Actually and originally, it did not do this *per se*, but rather, just wrote stdout directly to the `cwd()`. In Pulp, these operations were adopted *pro forma*, albeit bent by the obvious need of automated path construction *vis-a-vis* production environment standards.

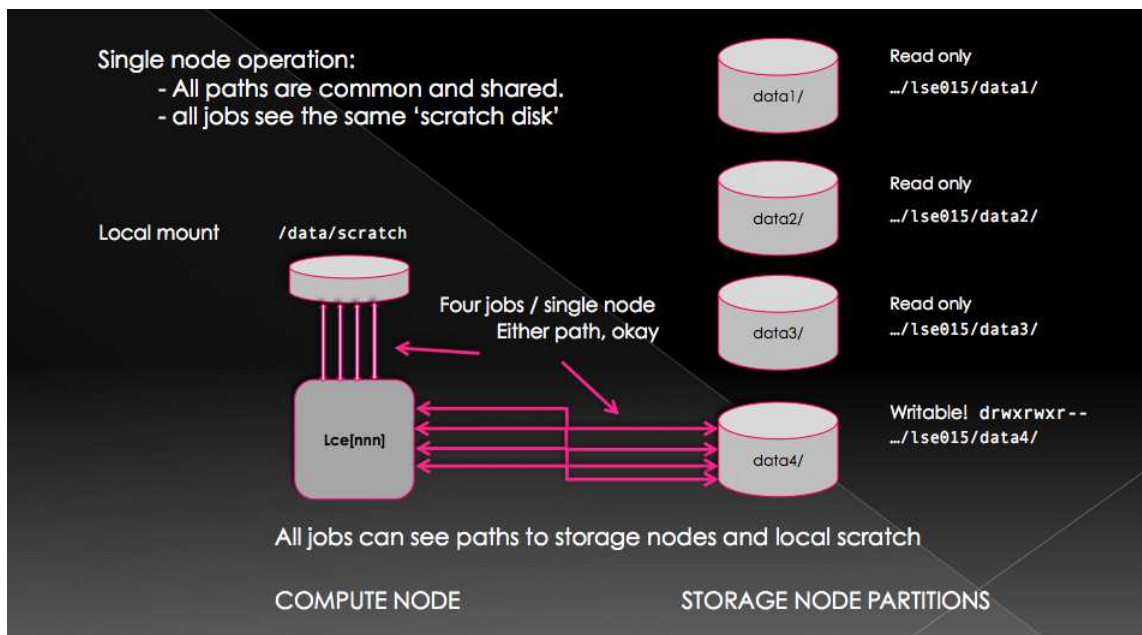


Figure 1: Single node operations: all paths are common and shared

The *ad hoc* addition of command line options, and specifically that of “--outputpath”, was advised against for the Pulp package. The driver to have a full and open output path appears to have been a desire to write to local disk. User’s must understand that an output path to local disk would *break* the current Pulp pipeline, unless a user were to restrict pipeline deployment to one (1) compute node.

Figure 1 demonstrates that, for a single node – as the pulsar shell script would operate on – all paths to all write devices are the same for each and every job the pipeline is running. This is certainly *not* true for framework pipeline operations on a multi-node cluster.

The reasons are illustrated in the figure 1. What the figure illustrates can be stated succinctly: when operating in a multi-node cluster compute environment, locally mounted node disks do not share common pathnames. Writing to local disk would require a complete redesign and would tend toward the programmatic operations of the standard imaging pipeline.

But even if Pulp were refactored to write to local disk. This still won’t work for the pulsar pipeline as it currently operates. Speaking specifically of the “RSPA” processing as it is done now, the current implementation of using symbolic links to “link in” all beam formed data files for “RSPA” will not work. Unless someone knows a way to create symbolic links across or throughout cluster compute nodes, “RSPA” processing cannot be done in a cluster environment. Restricting the whole pipeline to a single node, which is

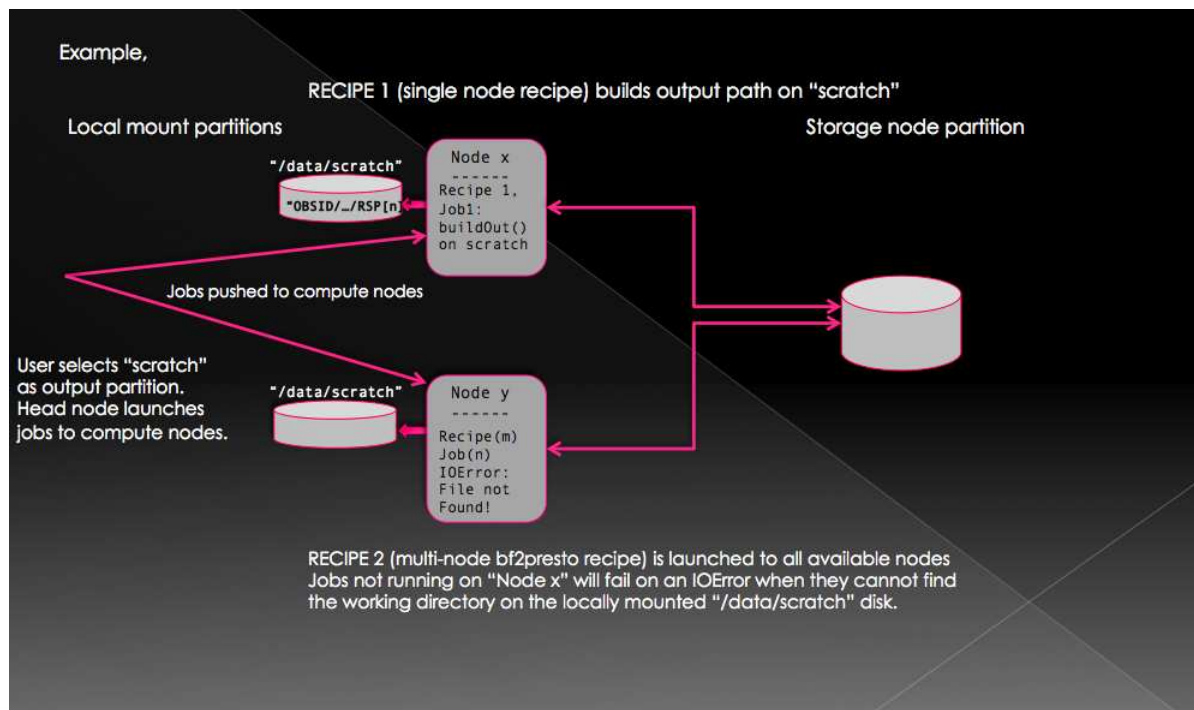


Figure 2: Writing to local disk in a multi-node environment. Paths are common, but not shared.

easily doable, would allow "RSPA" processing to go forward, but then that's not exactly cluster computing, is it?

## 4.5 Future enhancements

## 5 Acknowledgement

The author wishes to thank John Swinbank for providing vital insight and levity when the author had none.

## 6 Appendix A

### 6.1 Pulp Class Hierarchy: a comprehensive listing of public and private methods

The follow images are a schematic representation of Pulp class hierarchy. First, the top level of the package, showing the pipeline definitions, version and package files. (images: `idle` Path browser)

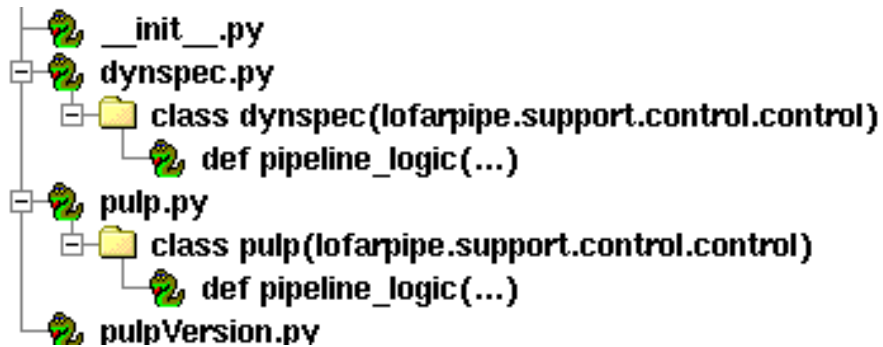


Figure 3: Pulp, v1.0, Class Structure and Methods and Functions, “definition” modules

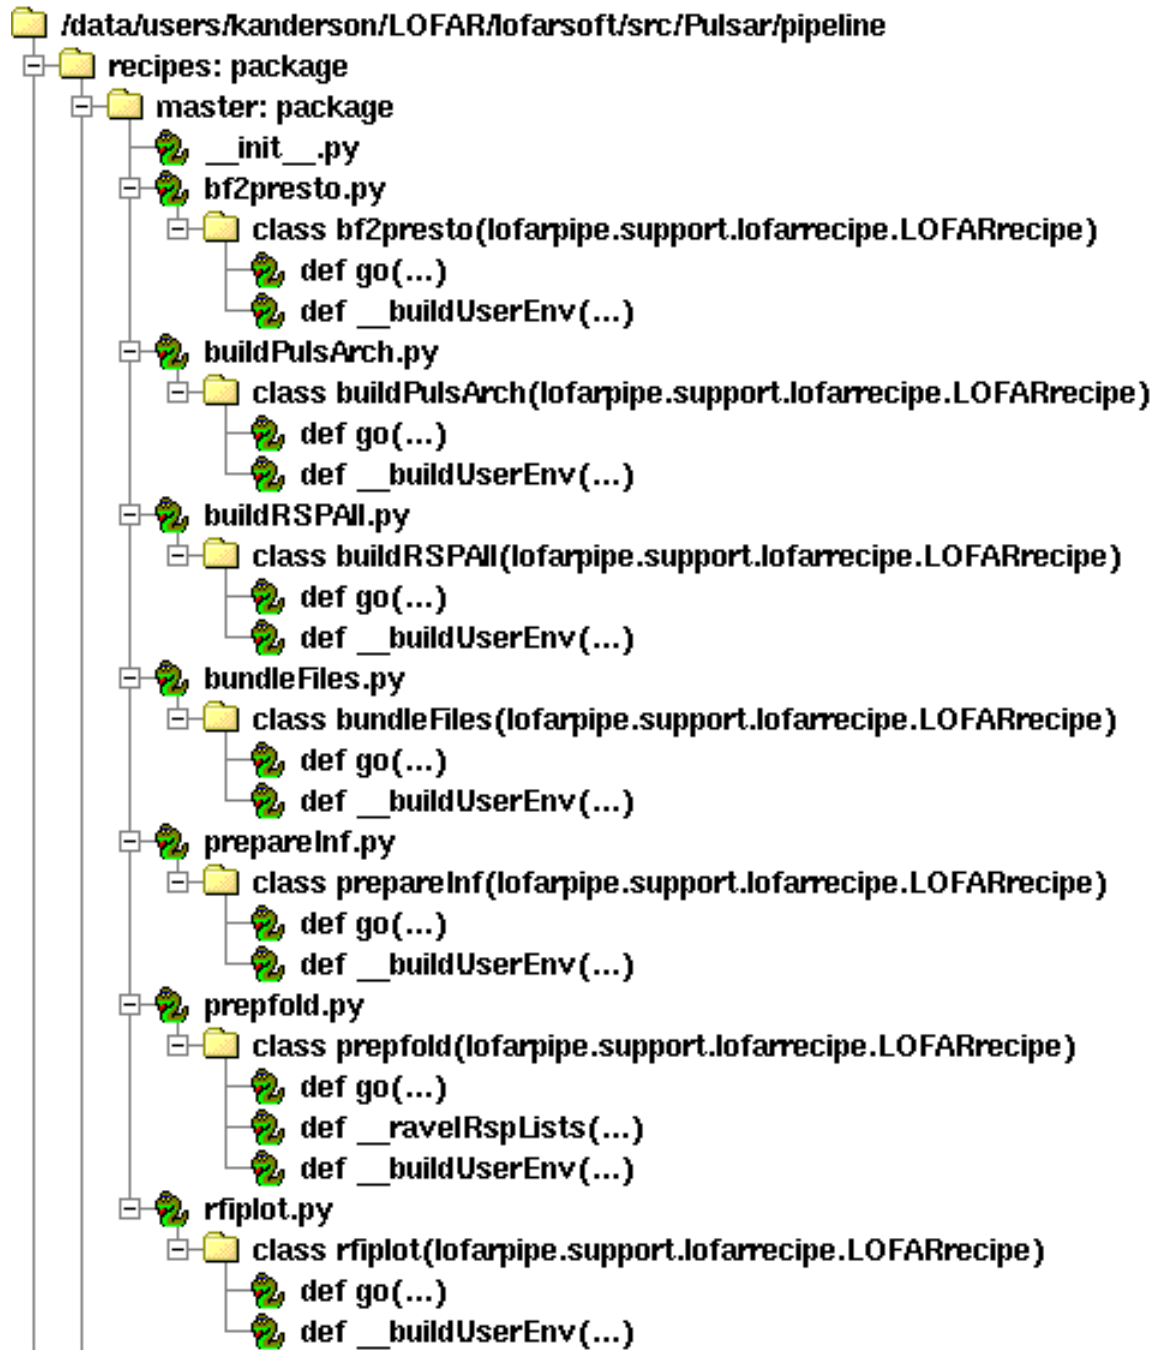


Figure 4: Pulp, v1.0, Class Structure, Methods and Functions, 'master' recipe modules




Figure 5: Pulp, v1.0, Hierarchical Class Structure and Methods, “node” modules



Figure 6: Pulp, v1.0, Hierarchical Class Structure and Methods, “support” modules